

Openstack-orbit

A Virtual Openstack Environment

Pedro Pinto da Silva
Prof. Luís Lopes

Contents

- Part 1** - Introduction and motivation
- Part 2** - Openstack Overview
- Part 3** - Prototype setup
- Part 4** - Validating and using Openstack

Part 1

Introduction and Motivation

What is Openstack?



Openstack is..

Open-source software for deploying and managing cloud computing platforms.

- Serves primarily as a solution of Infrastructure-as-a-Service (IaaS).

Who drives Openstack?

- Created in 2010 by **NASA** and **Rackspace Hosting**.



- The project is currently managed by Openstack Foundation, a non-profit organization (since 2012).



Releases



- **11 releases** so far
- **Current Release Cycle** - Every 6 months

Problem

1. Abundant but **fragmented** documentation
2. Existing tutorials **assume knowledge** by reader/viewer
3. How to create a basic **proof-of-concept** environment

Motivation

1. Explore Openstack - develop “**know-how**”
2. Install and configure a **basic** Openstack **setup**
 - a. Use a **community** driven package (**RDO**)

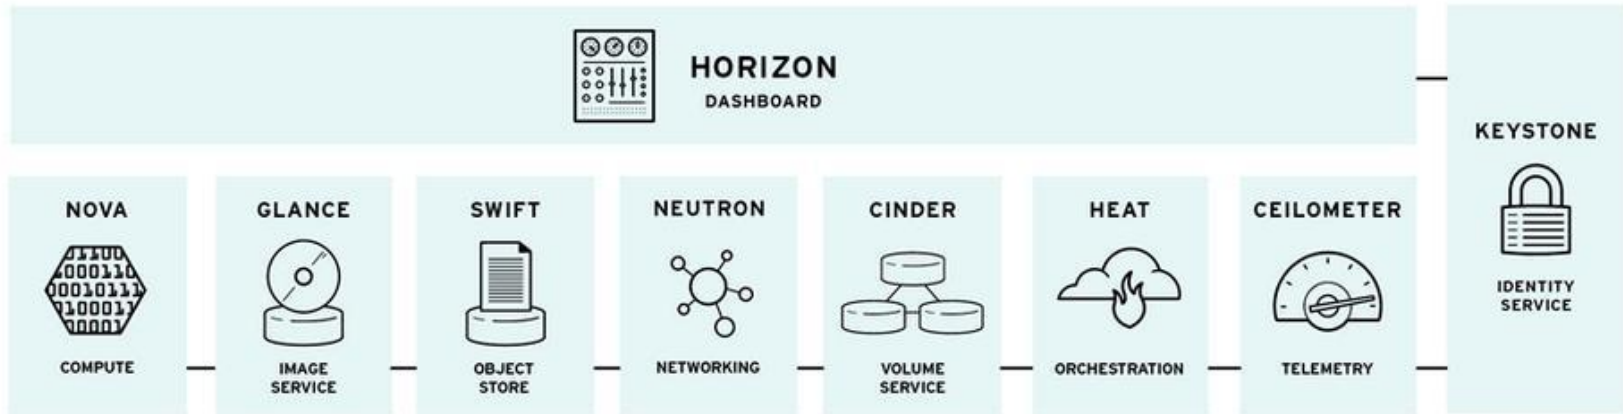


Part 2

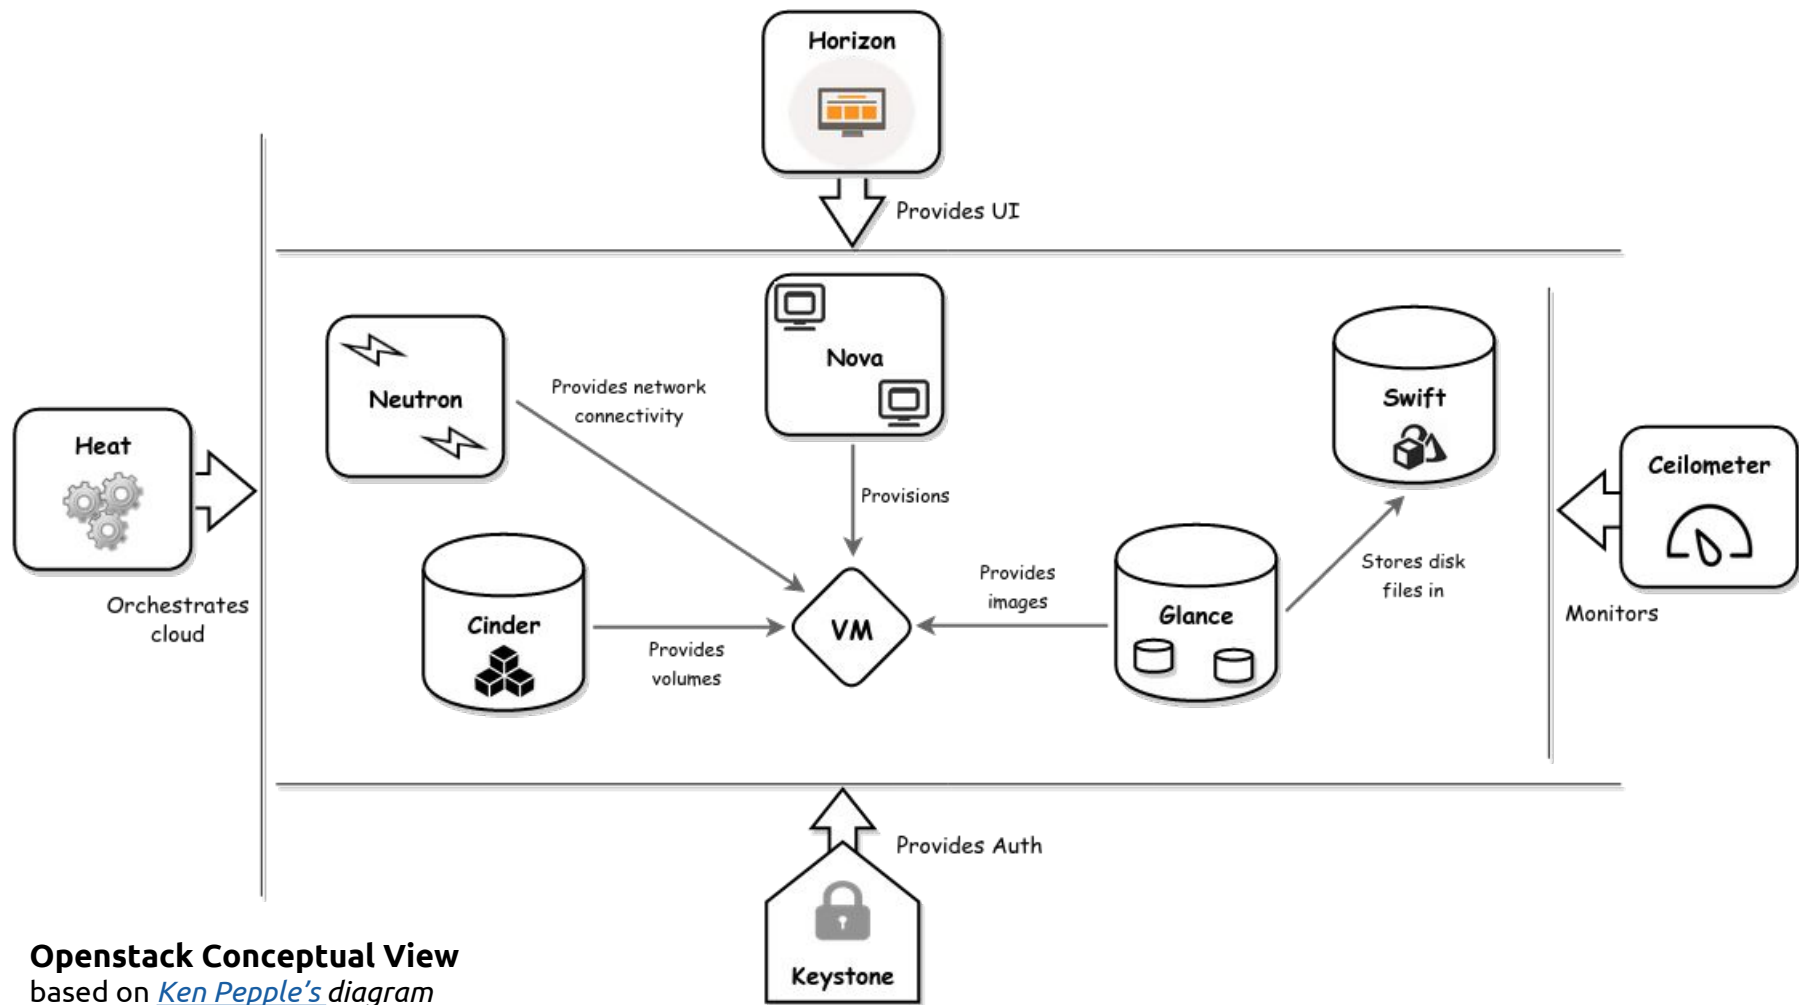
Openstack Overview

Openstack Services

- Openstack is divided into **different components**.
- **Integration** is achieved through **APIs** - offered and consumed by each service.




Openstack Services - by [RedHat](#)



Openstack Conceptual View
based on [Ken Pepple's](#) diagram

How are these components organized?



Basic Neutron Architecture

Controller:

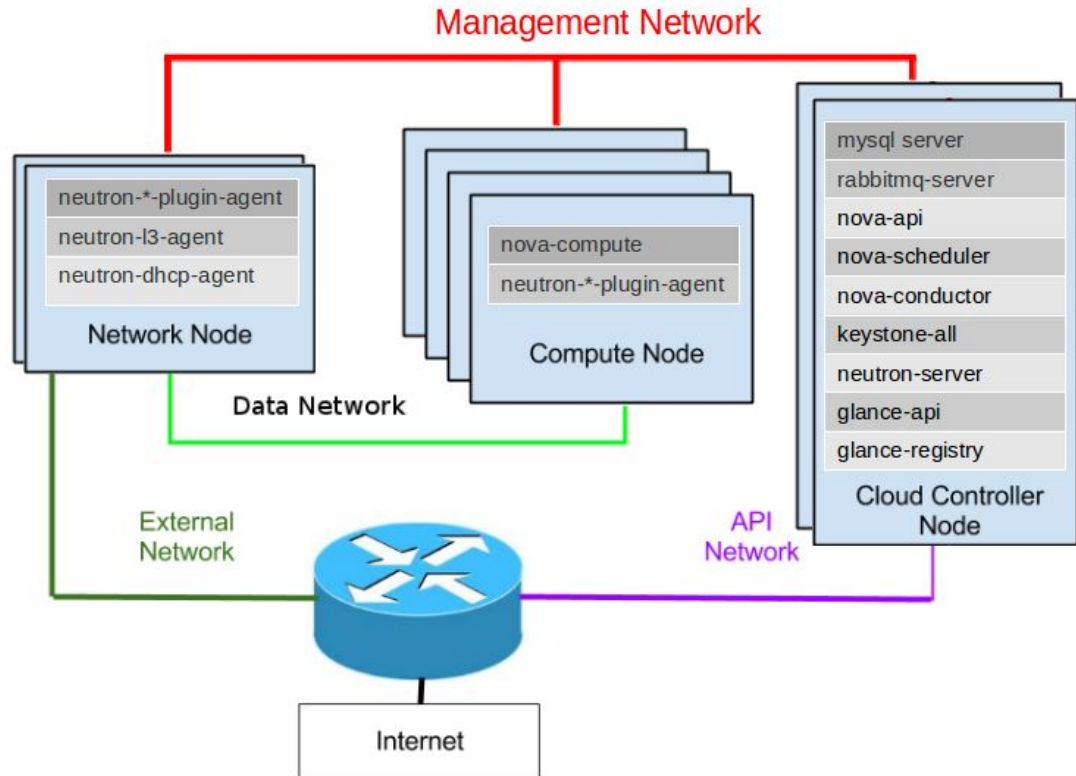
Runs **management** (identity, image) and **supporting** services (DB, AMQ, NTP).

Compute:

Runs the **hypervisor** that operates vm instances.

Network:

Runs the **networking plug-in** and neutron agents. Handles **internet connectivity** for vm instances.



Basic Neutron Architecture

Physical Networks

Management

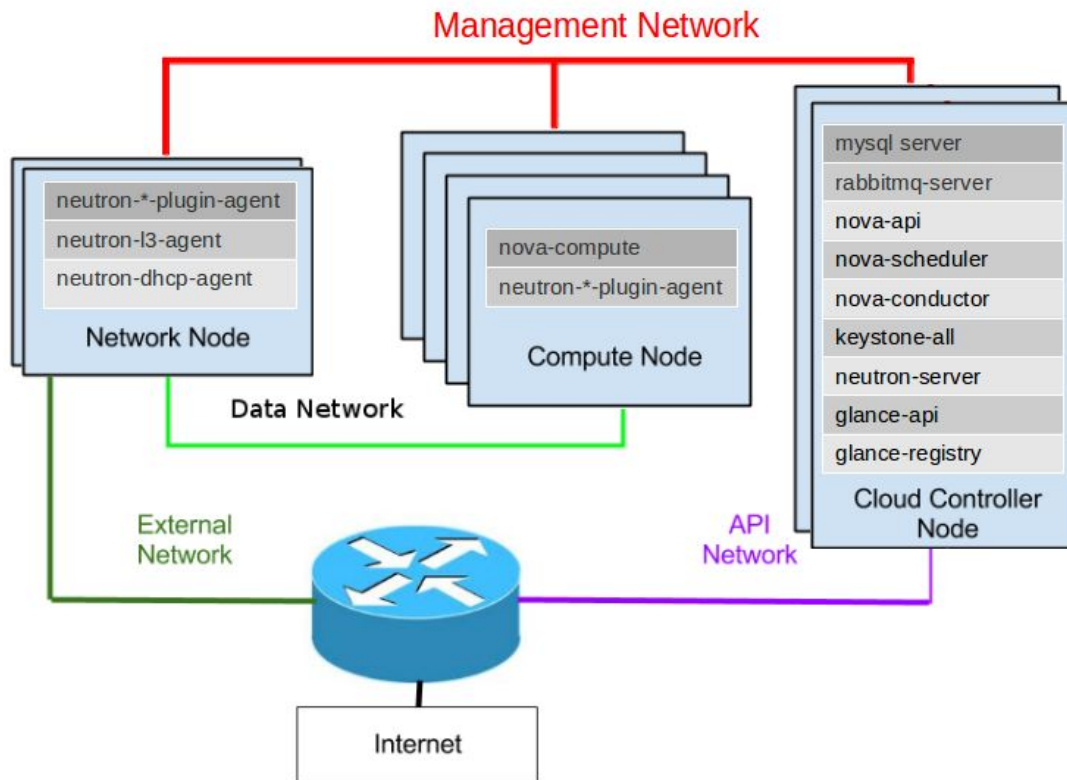
Provides **internal communication** between components.

Data

Provides **VM data communication** within the cloud deployment.

External/API

Exposes **VMs** and/or **Services** to consumers/users **outside** of Openstack.



Part 3

Prototype Setup

Setup Requirements

For a **testing environment** we need:

- 1 machine (min. 8 GB RAM)
- 1 hypervisor (KVM)
- 1 (minimal) installation image (RHEL-derived)
- 1 Openstack installation method (RDO-packstack)

Tools:

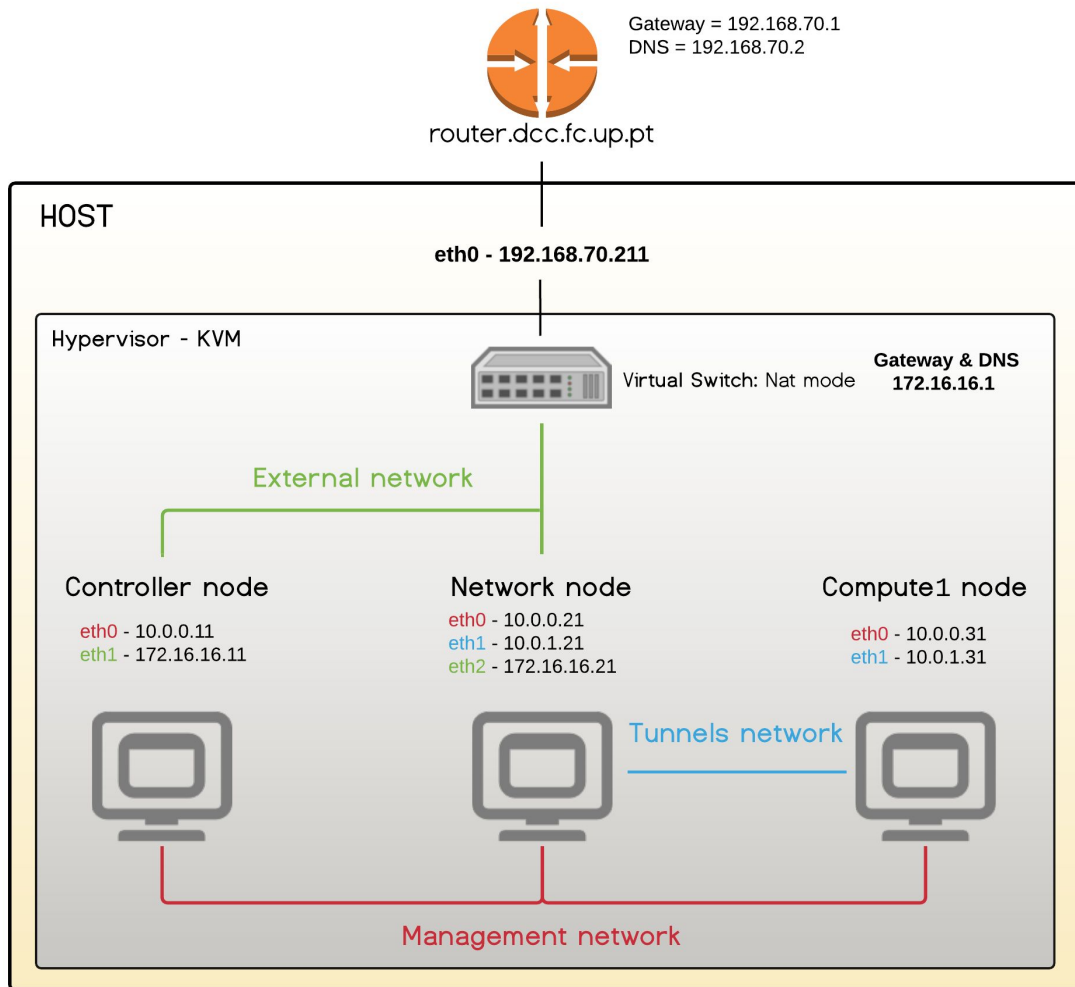
OS: CentOS 7.1
Release: RDO Kilo
Installer: Packstack
Hypervisor: KVM
Neutron Plugin: OpenVSwitch
Network Type: Gre-Tunnels



Options available:

flat, flat-dhcp
vlan
vxlan, gre

Prototype Setup
with Basic Neutron Architecture



RDO

Community supported distribution of Openstack on RHEL-derived distributions (CentOS, Fedora, Scientific Linux).

Why use RDO?

Upstream code specifically **packaged** and **patched** to run well on **CentOS**.



Packstack

An **installation utility** for **automating** the deployment of simple/small Openstack clouds with RDO (uses Puppet).

Requires:

- **Pre-installed** and **network-ready** servers.
- **Compatible OS** - Fedora and RHEL derivatives.

Packstack

Packstack **is** for:

- Single host or multi-node deployments (small-scale)
- Proof of Concept deployments (mostly)

Packstack **is not**:

- A server provisioning tool (systems must be set up in advance)
- Does not set up HA, Load-Balancing and Scalability options available in Openstack

Packstack answers-file

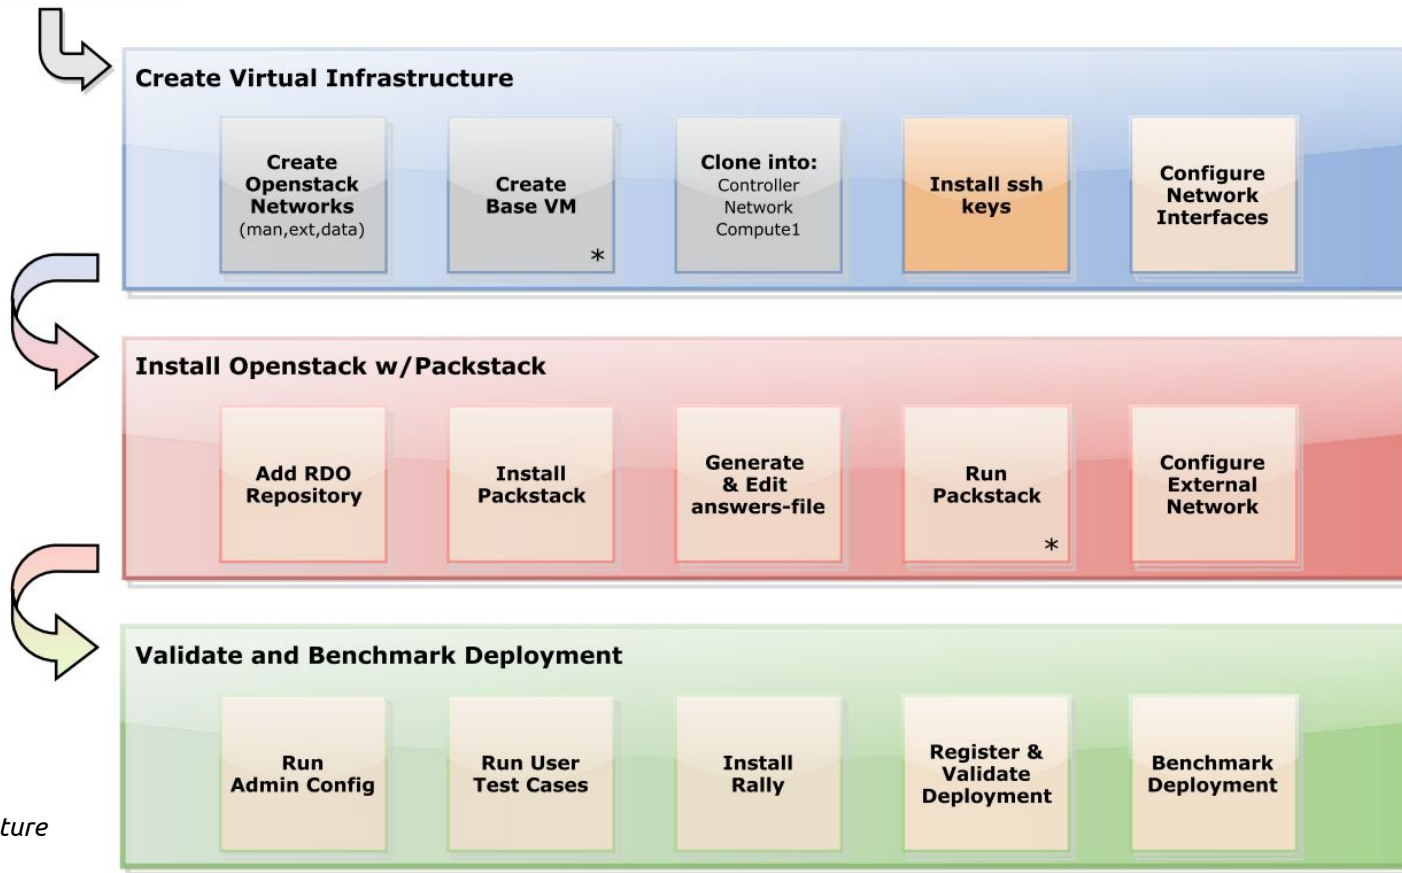
Ours differs from the default like this:

```
CONFIG_COMPUTE_HOSTS           = 10.0.0.31
CONFIG_NETWORK_HOSTS           = 10.0.0.21
CONFIG_NEUTRON_ML2_TENANT_NETWORK_TYPES = gre
CONFIG_NEUTRON_ML2_TYPE_DRIVERS  = gre
CONFIG_NEUTRON_ML2_TUNNEL_ID_RANGES = 1001:2000
CONFIG_NEUTRON_ML2_VNI_RANGES    = 1001:2000
CONFIG_NEUTRON_OVS_BRIDGE_IFACES = br-eth1:eth1
CONFIG_NEUTRON_OVS_BRIDGE_MAPPINGS = physnet1:br-eth1
CONFIG_NEUTRON_OVS_TUNNEL_IF     = eth1
CONFIG_PROVISION_DEMO            = n
```

Important!!

Default neutron plugin - OpenVSwitch

Initialization



Orbit Setup Workflow
with Basic Neutron Architecture

* - Time consuming tasks


Tasks performed on

Hypervisor

Guest

Both

Can we automate this procedure?



Orbit - A virtual environment

Motivation:

Automate the creation of a “PoC” (Proof-of-Concept) setup.

Features:

- Fully **automated** (conflicts resolved through y/n questions).
- Reads **configurations** from **file**.
- **Logs** to file and dumps verbose to stdout.
- **Tests** and **benchmarks** the environment (rally).

Orbit - Resulting System

Upon successful installation:

- Ready-to-use and tested basic Openstack setup (3-node).
- Custom benchmarking (edit templates/rally-tasks).

Considerations:

- Still a bit buggy:
 - Many points of failure
 - Resume installation from last execution point would be great

Part 4

Validating and using Openstack

Using Openstack

What a successful installation looks like:

```
10.3.0.21_postscript.pp: [ DONE ]
10.3.0.31_postscript.pp: [ DONE ]
10.3.0.11_postscript.pp: [ DONE ]
Applying Puppet manifests [ DONE ]
Finalizing [ DONE ]

**** Installation completed successfully ****

Additional information:
* Time synchronization installation was skipped. Please note that unsynchronized time on server instances might be
  problem for some OpenStack components.
* File /root/keystonerc_admin has been created on OpenStack client host 10.3.0.11. To use the command line tools y
  ou need to source the file.
* To access the OpenStack Dashboard browse to http://10.3.0.11/dashboard .
Please, find your login credentials stored in the keystonerc_admin in your home directory.
* To use Nagios, browse to http://10.3.0.11/nagios username: nagiosadmin, password: 495f9dfc4de94688
* The installation log file is available at: /var/tmp/packstack/20150914-085933-DNLra5/openstack-setup.log
* The generated manifests are available at: /var/tmp/packstack/20150914-085933-DNLra5/manifests
```

Ok.. **So what's next?**

Using Openstack

Create:

→ Tenants & Users

→ Networks

Private Networks - created by tenants

External Networks - created by admin


→ Routers

Connect networks (e.g. private to external)

→ Images, Flavours

→ Security rules, keys

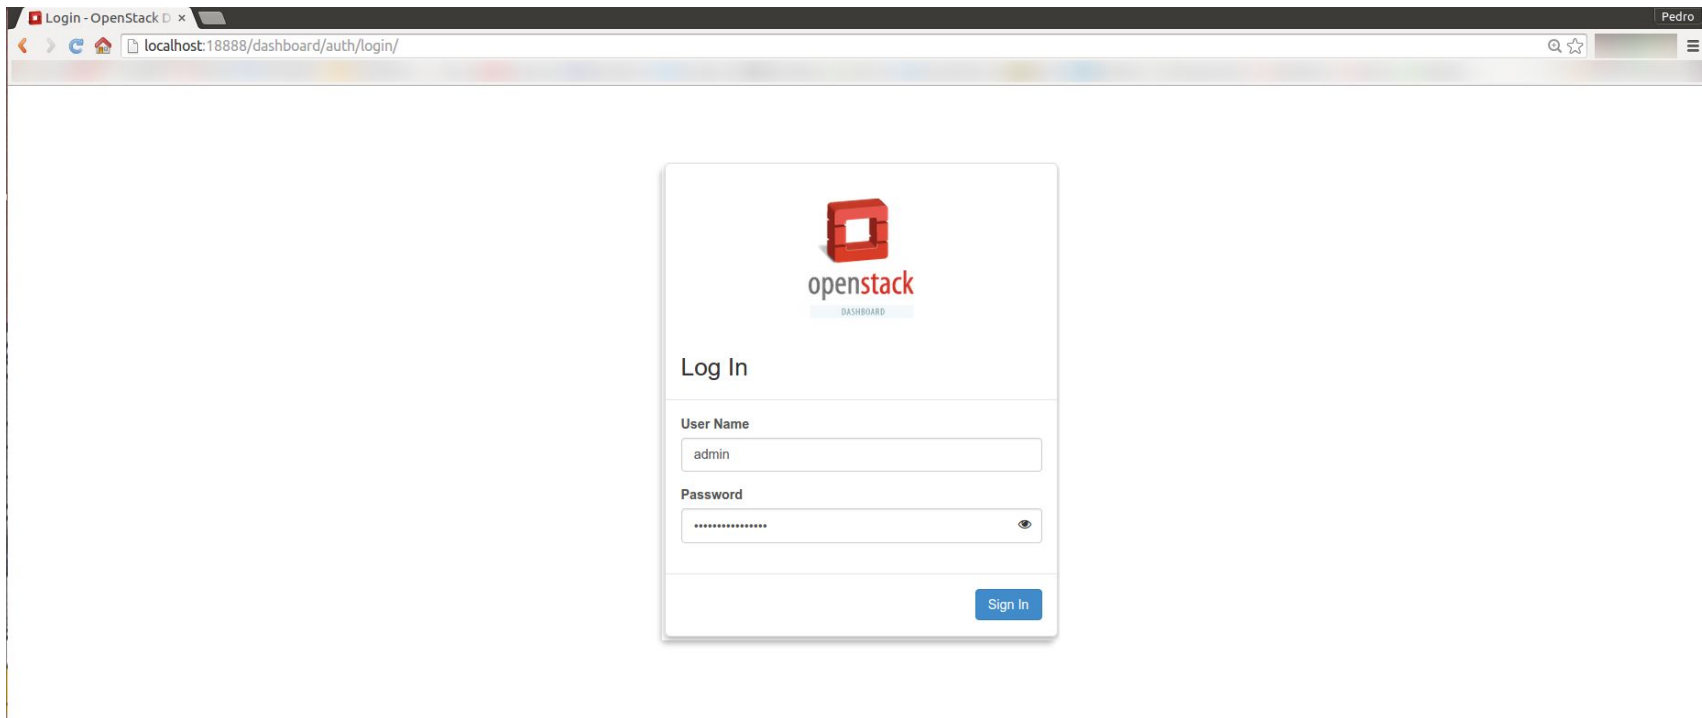
→ VM Instances



Use the dashboard / the cli (command-line interface)

Dashboard

SSH tunnel: `ssh -f -N -T -L 18888:10.3.0.11:80 psilva@192.168.70.211 -i ~/.ssh/id_rsa`



The screenshot shows a web browser window with the title "Login - OpenStack". The address bar displays "localhost:18888/dashboard/auth/login/". The page content features the OpenStack logo (a red cube) and the text "openstack DASHBOARD". Below this is a "Log In" section with two input fields: "User Name" containing the text "admin" and "Password" containing masked characters. A blue "Sign In" button is located at the bottom right of the login form.

Network Topology - x

localhost:18888/dashboard/project/network_topology/

admin

Network Topology

Small Normal

Launch Instance + Create Network + Create Router

Network Topology

Networks

Routers

Object Store

Admin

Identity

The diagram illustrates a network topology. On the left, a blue vertical line represents the 'external_network' with IP address 172.16.2.0/28. A router icon labeled 'router_ex..' connects this to an orange vertical line representing the 'private_network' with IP address 192.168.20.0/25. The router's interface on the external network is 172.16.2.2, and its interface on the private network is 192.168.20.1. A server icon labeled 'test_serv..' is connected to the private network at IP 192.168.20.3. A tooltip for the 'test_server' instance shows ID 3035f4a6-be00-4d5a-969a-2e140539e1f9, STATUS ACTIVE, and buttons for 'View Instance Details', 'Open Console', and 'Terminate Instance'.

external_network 172.16.2.0/28

router_ex.. Router

172.16.2.2

192.168.20.1

private_network 192.168.20.0/25

test_serv.. Instance

192.168.20.3

test_server
ID 3035f4a6-be00-4d5a-969a-2e140539e1f9
STATUS ACTIVE
» View Instance Details » Open Console Terminate Instance


Our test topology:

Very simple.

Why?

Validate that our instances have connectivity.

Network Topology - Example



Example Network Topology
by Rackspace

Deployment Validation

Created VM instances should:

1. Be accessible from the outside (icmp, ssh) - **floating IP**
2. Have access to the internet.
3. Be able to talk to each other directly and privately if on the same subnet (using the data network, over GRE tunnels)

Deployment Validation

```
# - Trying again in 15 seconds
# - Counter = 1
# - Instance is ACTIVE
# - Pinging instance..
PING 192.168.20.3 (192.168.20.3) 56(84) bytes of data.
64 bytes from 192.168.20.3: icmp_seq=1 ttl=64 time=7.41 ms
64 bytes from 192.168.20.3: icmp_seq=2 ttl=64 time=12.5 ms
64 bytes from 192.168.20.3: icmp_seq=3 ttl=64 time=11.8 ms
64 bytes from 192.168.20.3: icmp_seq=4 ttl=64 time=14.8 ms

--- 192.168.20.3 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3005ms
rtt min/avg/max/mdev = 7.411/11.669/14.860/2.701 ms
[ OK ]
[orbit] Create a floating ip for inbound access to test_server..
+-----+-----+-----+-----+-----+
| Id | IP | Server Id | Fixed IP | Pool |
+-----+-----+-----+-----+-----+
| f406691d-9dd6-44f2-9a91-a40b5e7eba3 | 172.16.2.3 | - | - | external_network |
+-----+-----+-----+-----+-----+
[ OK ]
[orbit] Associate the floating ip with the test server instance..
[ OK ]
[orbit] Configure network node to allow external traffic coming from and into the instance..
iptables: Saving firewall rules to /etc/sysconfig/iptables: [ OK ]
[ OK ]
```

(a) Ping Instance through namespace

Ping instance through network namespace and create a floating IP.

```
[orbit] Check if the test server is reachable by its floating ip: with icmp..
PING 172.16.2.3 (172.16.2.3) 56(84) bytes of data.
64 bytes from 172.16.2.3: icmp_seq=1 ttl=63 time=21.8 ms
64 bytes from 172.16.2.3: icmp_seq=2 ttl=63 time=1.96 ms
64 bytes from 172.16.2.3: icmp_seq=3 ttl=63 time=1.44 ms
64 bytes from 172.16.2.3: icmp_seq=4 ttl=63 time=1.72 ms

--- 172.16.2.3 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3005ms
rtt min/avg/max/mdev = 1.442/6.751/21.872/8.732 ms
[ OK ]
[orbit] Check if the test server is reachable by its floating ip: with ssh..
# 172.16.2.3 SSH-2.0-dropbear_2012.55
# 172.16.2.3 SSH-2.0-dropbear_2012.55
Im alive at last - cirros instance
Can I ping google?
PING 8.8.8.8 (8.8.8.8): 56 data bytes
64 bytes from 8.8.8.8: seq=0 ttl=52 time=11.574 ms
64 bytes from 8.8.8.8: seq=1 ttl=52 time=14.628 ms
64 bytes from 8.8.8.8: seq=2 ttl=52 time=8.315 ms
64 bytes from 8.8.8.8: seq=3 ttl=52 time=14.569 ms

--- 8.8.8.8 ping statistics ---
4 packets transmitted, 4 packets received, 0% packet loss
round-trip min/avg/max = 8.315/12.271/14.628 ms
[ OK ]
```

(b) Ping internet from instance

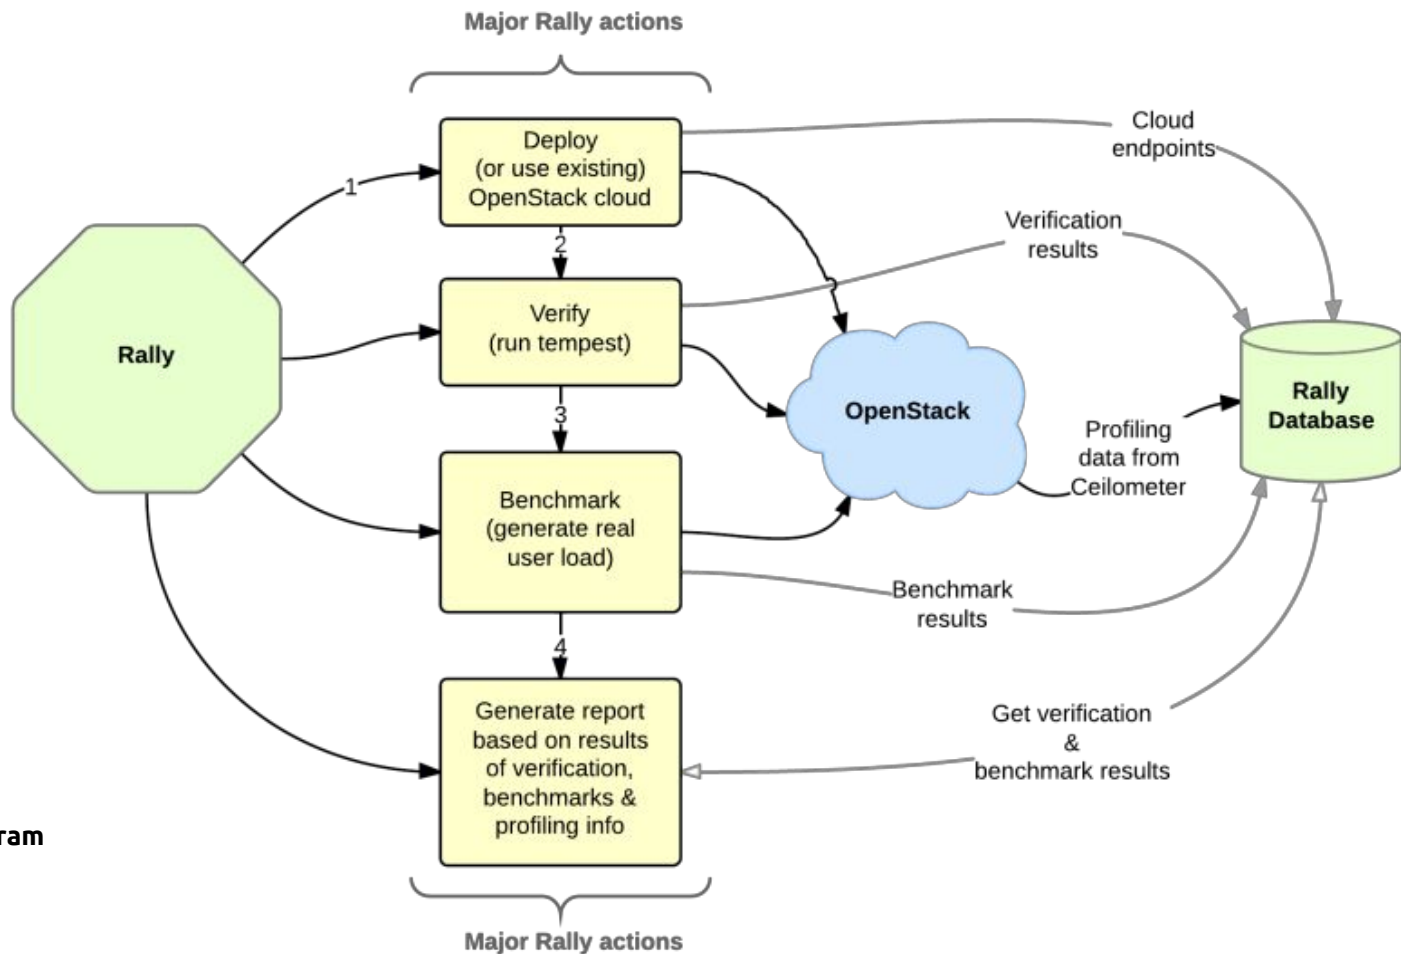
*Ping and ssh into instance through floating IP.
Ping 8.8.8.8 from within the instance (ssh).*

Openstack - Benchmarking

**How well does Openstack
work at scale?**

Are there any tools that simulate real work loads?

Rally



Rally Conceptual Diagram
by Rally

Rally - Tasks

```
tasks_dir="/usr/share/rally/samples/tasks/"  
  
"scenarios/"  
  
"authenticate/keystone.json"  
  
"vm/boot-runcommand-delete.json"  
  
"neutron/create_and_delete_networks.json"  
  
"nova/boot-and-live-migrate.json"  
  
"nova/boot-and-delete-multiple.json"  
  
"cinder/create-and-attach-volume.json"  
  
...
```

Tasks are defined in JSON

```
{% set flavor_name = flavor_name or "m1.tiny" %}  
{  
  "NovaServers.boot_and_delete_server": [  
    {  
      "args": {  
        "flavor": {  
          "name": "{{flavor_name}}"  
        },  
        "image": {  
          "name": "^cirros.*uec$"  
        },  
        "force_delete": false  
      },  
      "runner": {  
        "type": "constant",  
        "times": 10,  
        "concurrency": 2  
      },  
      "context": {  
        "users": {  
          "tenants": 3,  
          "users_per_tenant": 2  
        }  
      }  
    }  
  ]  
}
```

Example

Benchmark overview

Input file

▼ KeystoneBasic

create_delete_user

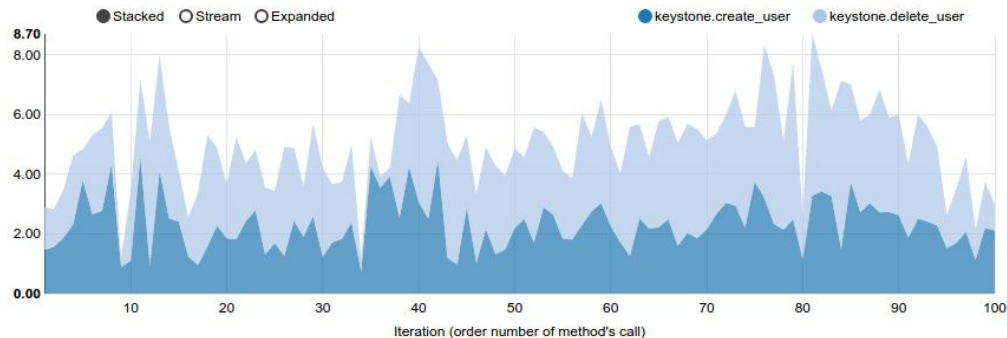
KeystoneBasic.create_delete_user (56.226s)

Overview

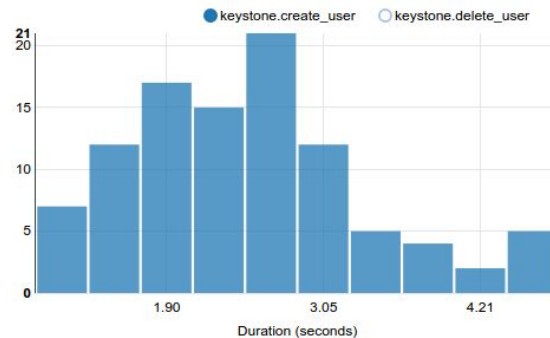

Details

Input task

Charts for each Atomic Action



● keystone.create_user ● keystone.delete_user



Square Root Choice ▼

Example of a task report
task: sla/create_delete_user

Orbit - Video demo

Under production!

Openstack in Production

Ubuntu:

MAAS + JUJU

RHEL:

OpenCrowbar + Packstack

Hardware provisioning, O/S
and networking configuration

Install software requirements
and Openstack

Tools for deploying Openstack at scale!

Discussion

To sum up:

- Very **powerful** but **complex** software.
- Covers a wide variety of **use cases**.
- Promotes **self-provisioning** (less stress into Sys-Admin).
- Possible to automate - to some extent, with the right tools.

To do

Choose between:

- **Enhance** orbit-tools: + usability & robustness.
- **Develop** orbit-tools: add/remove nodes (automatically + benchmark)
- Study **Storage** options (ephemeral, persistent, object, block)
- Other?

The End

Extra Material - For consultation

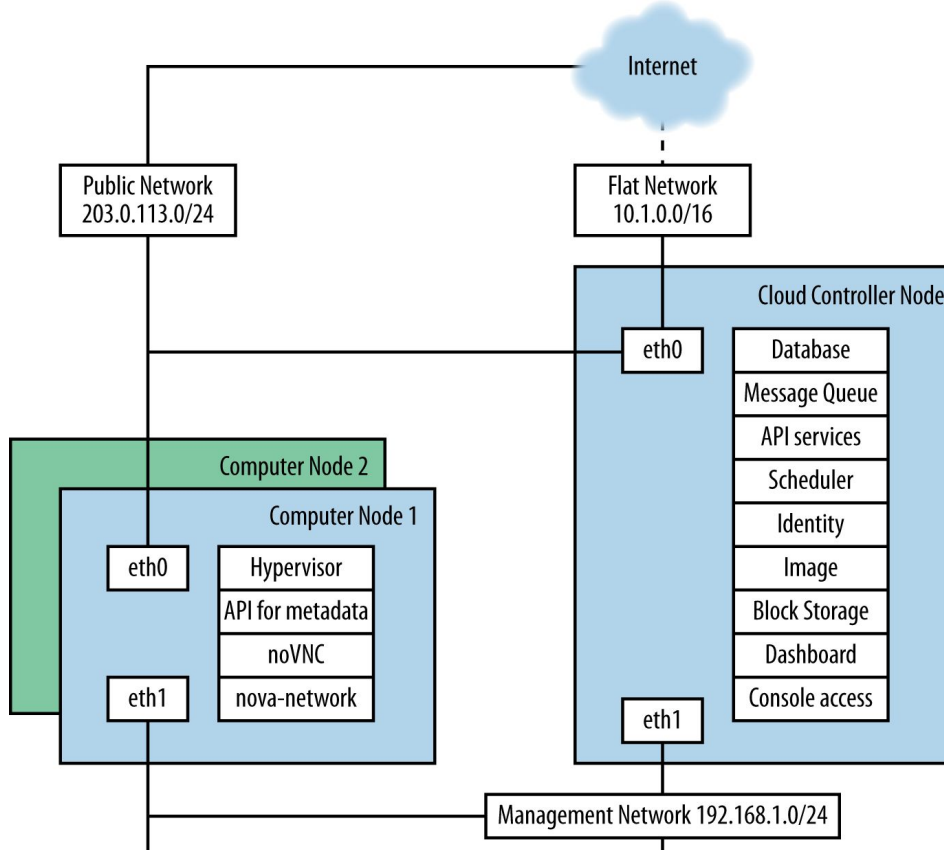
Before Neutron networking there is...

Nova-Network

Limitations:

- **Network topologies** (L2 only)
- **Scaling** (VLAN dhcp)
- No **3rd party** network solutions
- Requires **more** manual/hours-of **configuration**

Obsolete as soon as Neutron is stable
(deprecated)



Legacy Networking Architecture (nova) - by Openstack

Why Openstack Neutron

- Give tenants an **API** to build **L3 networks** (self-provisioning)
- Allow **3rd party** backend **plugins** (VMware NSX, Cisco UCS, OVS, ...)
- Advanced network services (LB-aaS, VPN-aaS, firewall-aaS, ...)
- Surpass VLAN limitations (build overlay networks using gre/vxlan)

Software Defined Networks

Neutron is better (more flexible) at software defined networking:

Separation of the control and data planes.

Control plane - logic for controlling forwarding behavior

Data plane - forward traffic according to logic

Current hardware solutions have both implemented and tightly coupled.

By separating these, software can grow independently of hardware. (and not be restricted by it - proprietary implementations)

Good for datacenters - VM migration, L2 Switching - guarantee consistency.

Then you can program the network switches from a central database/controller.

OpenVSwitch Plugin

OpenVSwitch is the default network plug-in for neutron.

- Performs ethernet switching in a hypervisor (beyond simply bridging)
- Implements Openflow (you can use any field in a frame (L2 to L4) and make decisions on them)
- Allows automation in virtualized network infrastructures
- Makes it possible to build sophisticated networks with not so sophisticated hardware
- Is capable of automating the kind of virtualized network infrastructure that a IaaS cloud provider needs

Gre-Tunnels

Gre - Generic Routing Encapsulation (**GRE**) is a tunneling protocol developed by Cisco Systems that can encapsulate a wide variety of network layer protocols inside virtual point-to-point links over an Internet Protocol network.

gre and **vxlan** networks are very similar.


They are both "overlay" networks that work by encapsulating network traffic.

Like **vlan** networks, each network you create receives a unique tunnel id.

Unlike **vlan** networks, an overlay network does not require that you synchronize your OpenStack configuration with your L2 switch configuration.


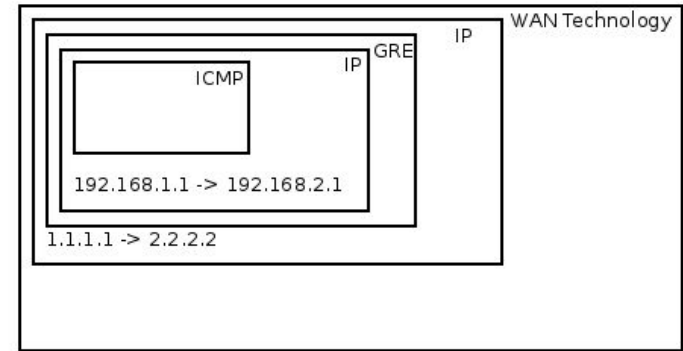
Gre-Tunnels 2

Example



The goal is to enable hosts in the Tel-Aviv branch the ability to ping hosts in the Brno branch. Without a tunnel this would not be possible. It's done via statically configuring a GRE tunnel on both site routers, and requires no involvement from the different ISPs. If encryption is desired it can be later configured on top of the GRE tunnel. Two things need to be done on each router:

1. Bringing up a (virtual) tunnel device
2. Defining static routing to the remote 192.168.x.0/24 network through the local tunnel



Gre-Tunnels 3

When utilizing a GRE tunnel you have routing at your finger tips. The forwarding decision over the tunnel is based on a routing decision. You can utilize all the dynamic routing protocols over the tunnel interface as if it were a physical interface. The GRE tunnel can be encrypted and when it's sent out over public networks such as the internet it should be encrypted. The encryption decision for a GRE tunnel is based on the source/destination addresses of the tunnel and not the traffic going through the tunnel. This simplifies the encryption decision into a 1 line ACL. It also decouples the forwarding decision from the encryption decision.

When utilizing a standard IPSEC tunnel the forwarding decision is based on an ACL. In essence the forwarding decision is based on the encryption decision. This decision is based on an ACL that is static and must be updated to increase or decrease functionality. This critical ACL must be an exact match at both locations to ensure stability.

Since I'm a WAN routing fellow I'm always more comfortable having a forwarding decision that's based on routes/routing protocols versus static ACL's.

GRE has a bit more overhead (24 bytes), but it's well worth it for the functionality gained.

Floating IP

A floating IP address is a service provided by Neutron. It's not using any DHCP service nor being set statically within the guest. As a matter of fact the guest's operating system has completely no idea that it was assigned a floating IP address. The delivery of packets to the interface with the assigned floating address is the responsibility of Neutron's L3 agent. Instances with assigned floating IP address can be accessed from the public network by the floating IP.

Floating IP address and a private IP address can be used at the same time on a single network-interface. The private IP address is likely to be used for accessing the instance by other instances in the private network while the floating IP address would be used for accessing the instance from a public network.

Source: https://www.rdoproject.org/Difference_between_Floating_IP_and_private_IP

Storage Back-ends

Table 6.2. Persistent file-based storage support

	Object	Block	File-level ^a
Swift	✓		
LVM		✓	
Ceph	✓	✓	Experimental
Gluster	✓	✓	✓
NFS		✓	✓
ZFS		✓	
Sheepdog	✓	✓	

^aThis list of open source file-level shared storage solutions is not exhaustive; other open source solutions exist (MooseFS). Your organization may already have deployed a file-level shared storage solution that you can use.

Orbit - Usage

Command Line Interface:

```
[psilva@cracs-cloud01 openstack-orbit]$ ./install_orbit.sh --help
Usage: install_orbit.sh [options]
  --clean [cfg-file]  Clean installation (remove all traces). Parameters specified in cfg-file
  --save-base-vm      Save base vm - used for cloning any virtual node
  --skip-base-vm      Use a saved base vm - with name specified in orbit.cfg
  --debug             Do not clean anything in case installation fails
  --help              Prompt usage and help information

[psilva@cracs-cloud01 openstack-orbit]$
```

Orbit - Configuration file

```
# You can change the variables here - used in install orbit.sh
# Beware when changing these variables - do not let them unset!
```

```
# Config NAME/ID
```

```
config_id="pedros"
```

```
# KVM uri
```

```
kvm_uri="qemu:///system"
```

```
# Path to disk images
```

```
img_disk_path="/var/lib/libvirt/images"
```

```
# VM user - used to communicate over ssh with the machine
```

```
vm_user="admin"
```

```
vm_pass="admin123"
```

```
vm_root_pass="naoseiroot"
```

```
# Networks
```

```
management_network_name="orbit-management"
```

```
management_network_ip="10.3.0.1"
```

```
management_bridge_name="virbr-man"
```

```
management_network_netmask="255.255.255.0"
```

```
data_network_name="orbit-tunnels"
```

```
data_network_ip="10.3.1.1"
```

```
data_bridge_name="virbr-tun"
```

```
data_network_netmask="255.255.255.0"
```

```
ext_network_name="orbit-external"
```

```
ext_network_ip="172.16.1.1"
```

```
ext_bridge_name="virbr-ext"
```

```
ext_network_ip_start="172.16.1.2"
```

```
ext_network_ip_end="172.16.1.128"
```

```
ext_network_netmask="255.255.255.0"
```

```
floating_network="172.16.2.1/28"
```

```
test_tenant_network="192.168.20.0/25"
```

```
# Args base vm
```

```
vm_base_name="orbit-base"
```

```
vm_base_size=12 #In GB
```

```
vm_base_ram=2048 #In MB
```

```
vm_base_vcpus=1
```

```
# Args Controller
```

```
vm_controller_name="orbit-controller"
```

```
vm_controller_ip_man="10.3.0.11" # Management network
```

```
vm_controller_ip_ext="172.16.1.11" # External/API network
```

```
# Args Compute1
```

```
vm_computel_name="orbit-computel"
```

```
vm_computel_ip_man="10.3.0.31" # Management network
```

```
vm_computel_ip_tun="10.3.1.31" # Data network
```

```
vm_computel_ip_ext="172.16.1.31" # Dummy
```

```
#Args Network
```

```
vm_network_name="orbit-network"
```

```
vm_network_ip_man="10.3.0.21" # Management network
```

```
vm_network_ip_tun="10.3.1.21" # Data network
```

```
vm_network_ip_ext="172.16.1.21" # External network
```

```
# SSH Key Name
```

```
ssh_key_name="orbit-key"
```

```
# NTP servers
```

```
ntp_servers_list="ntp01.fccn.pt,ntp02.fccn.pt"
```

Orbit - Tree

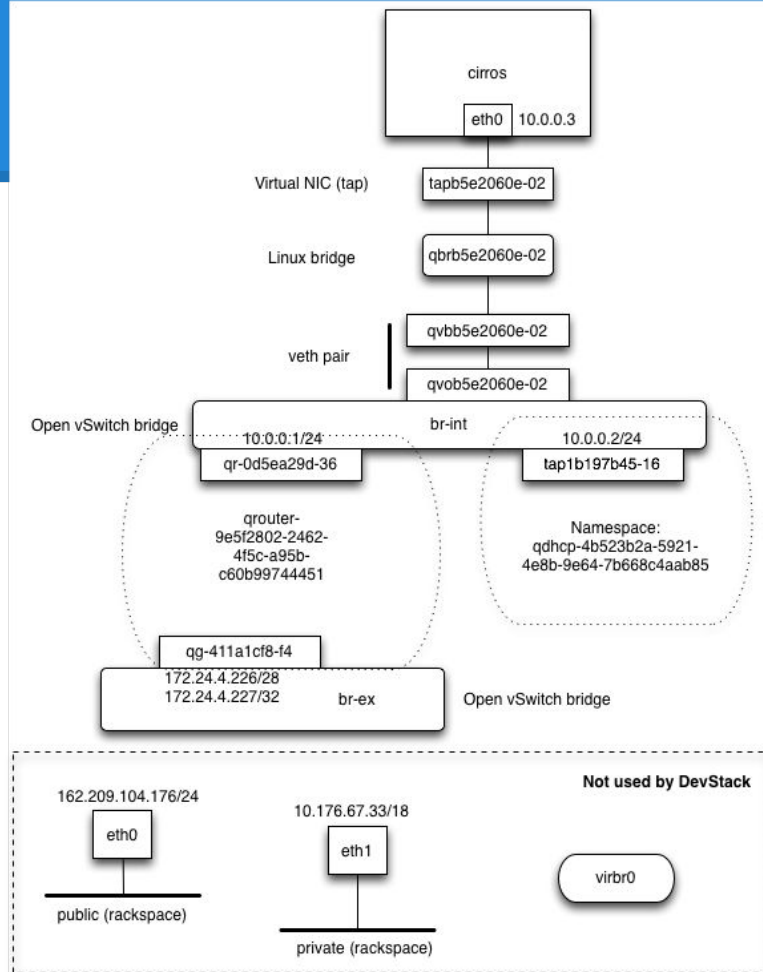

```
[psilva@cracs-cloud01 openstack-orbit]$ tree
```

```
├── functions
│   ├── add-net-interface
│   ├── clone-vm
│   ├── create-vm-ks
│   ├── delete-kvm-network
│   ├── delete-vm
│   ├── macgen-kvm
│   └── remove-net-interface
├── install_orbit.sh
├── LICENSE
├── orbit.conf
├── orbit.log
├── orbit.pedros
├── README.md
├── scripts
│   ├── config-iface.sh
│   ├── config-ovs-bridge.sh
│   ├── config-ovs-port.sh
│   ├── reorder-ifaces.sh
│   └── set-ntp-dcc.sh
└── templates
    ├── isolated-network.xml
    ├── nat-network.xml
    ├── orbit-centos7.ks
    └── rally-tasks
```

```
3 directories, 22 files
```

```
[psilva@cracs-cloud01 openstack-orbit]$ ls -la
total 108
drwxrwxr-x 6 psilva psilva 4096 Set 14 16:25 .
drwx----- 10 psilva psilva 4096 Set 14 16:25 ..
drwxrwxr-x 2 psilva psilva 4096 Jul 15 18:20 functions
drwxrwxr-x 8 psilva psilva 4096 Set 10 15:38 .git
-rwxr-xr-x 1 psilva psilva 49093 Set 14 16:24 install_orbit.sh
-rw-r--r-- 1 psilva psilva 16384 Set 14 16:26 install_orbit.sh.swp
-rw-rw-r-- 1 psilva psilva 1088 Jun 8 18:26 LICENSE
-rw----- 1 psilva psilva 1705 Set 14 16:16 orbit.conf
-rw-rw-r-- 1 psilva psilva 263 Set 14 16:24 orbit.log
-rw-rw-r-- 1 psilva psilva 1712 Set 14 15:01 orbit.pedros
-rw-rw-r-- 1 psilva psilva 190 Jun 8 18:30 README.md
drwxrwxr-x 2 psilva psilva 4096 Jul 17 00:15 scripts
drwxrwxr-x 2 psilva psilva 4096 Set 8 14:59 templates
[psilva@cracs-cloud01 openstack-orbit]$
```

Network Traffic



Rally SLA

Benchmark overview

Input file

▼ KeystoneBasic

create_delete_user

KeystoneBasic.create_delete_user (56.226s)

Overview Details Input task

Load duration: 51.746 s Full duration: 56.226 s Iterations: 100 Failures: 0

Service-level agreement

Criterion	Detail	Success
max_seconds_per_iteration	Maximum seconds per iteration 8.69s <= 4.00s - Failed	False
failure_rate	Failure rate criteria 0.00% <= 0.00% <= 1.00% - Passed	True
outliers	Maximum number of outliers 0 <= 1 - Passed	True
max_avg_duration	Average duration of one iteration 5.06s <= 3.00s - Failed	False

Total durations

Action	Min (sec)	Median (sec)	90%ile (sec)	95%ile (sec)	Max (sec)	Avg (sec)	Success	Count
keystone.create_user	0.738	2.278	3.567	4.107	4.597	2.316	100.0%	100
keystone.delete_user	0.305	2.652	4.117	5.129	5.622	2.741	100.0%	100
total	1.188	5.062	7.121	7.709	8.694	5.058	100.0%	100

Charts for the Total durations

